

Streaming in Large Language Models (LLMs)

What is Streaming?

In Large Language Models (LLMs), **Streaming** refers to the process where the model begins sending output tokens (words or phrases) **as soon as they are generated**, rather than waiting for the entire response to be completed.

This mechanism allows users to view the model's response in real-time, similar to how humans type or speak—creating a more natural and interactive conversational experience.

Why Streaming is Important

1. Faster Response Time

The user starts seeing output immediately, reducing perceived latency and **drop-off rates**.

2. Human-like Interaction

Mimics **human conversation flow**, building trust and keeping users engaged.

3. Crucial for Multi-modal Interfaces

Enables dynamic UI updates, essential for chatbots with text, images, or voice.

4. Better UX for Long Outputs

Especially useful for **long code or text generation**, where users can follow along.

5. Token Efficiency

Users can **cancel mid-generation**, saving compute and token usage.

6. UI Interleaving Possibilities

Developers can display *“Thinking...”*, show tool outputs, or interleave messages dynamically.

How Streaming Works in LLMs

During text generation, the LLM produces output tokens sequentially. Streaming APIs deliver these tokens incrementally through a continuous event stream (e.g., via WebSockets or HTTP chunks).

This allows the frontend (like a chat UI) to render tokens as they arrive, achieving smooth, real-time updates.

Implementing Streaming in Streamlit

Basic Concept

Streamlit allows dynamic updates of text areas or containers as data arrives. The idea is to use a loop that listens for new tokens and updates the UI progressively.

Example Implementation

Streamlit Code Example

```
import streamlit as st
from langgraph_backend import chatbot
from langchain_core.messages import HumanMessage
CONFIG = {'configurable': {'thread_id': 'thread-1'}}
if "message_history" not in st.session_state:
    st.session_state.message_history = []
message_history = st.session_state.message_history
for message in message_history:
    with st.chat_message(message["role"]):
        st.markdown(message["content"])
user_input = st.chat_input("Ask me anything...")
if user_input:
    st.session_state.message_history.append({'role': 'user', 'content':
    user_input})
    with st.chat_message("user"):
        st.markdown(user_input)
    with st.chat_message("assistant"):
        with st.spinner("Thinking..."):
            ai_message = st.write_stream(
                message_chunk.content for message_chunk, metadata in
                chatbot.stream(
                    {'messages': [HumanMessage(content=user_input)]},
                    config=CONFIG,
                    stream_mode='messages'
                )
            )
        st.session_state.message_history.append({"role": "assistant",
        "content": ai_message})
```

Explanation

- `st.empty()` creates a placeholder for live updates.
- The `with client.chat.completions.stream(...)` block streams tokens.
- Each token is displayed as it arrives, creating a **real-time output effect**.
- Once complete, the final output replaces the temporary cursor (`|`).

Benefits in Practice

Streaming in Streamlit enables conversational UIs that feel instant, natural, and responsive. It is especially valuable for applications like:

- AI Chat Assistants
- Code Generation Tools
- Research Copilots
- Real-time Data Summarizers

Conclusion

Streaming transforms static AI responses into **dynamic, human-like interactions**. It enhances usability, improves perception of intelligence, and creates a seamless user experience—key to next-generation AI applications.